

Conversion Constructors

Introduction

When designing classes, you might find that certain data types can logically be converted into objects of the type you're creating. For example, when designing the C++ `string` class, you might note that `char *` C strings could have a defined conversion to `string` objects. In these situations, it may be useful to define *implicit conversions* between the two types. To define implicit conversions, C++ uses *conversion constructors*, constructors that accept a single parameter and initialize an object to be a copy of that parameter.

While useful, conversion constructors have several major idiosyncrasies, especially when C++ interprets normal constructors as conversion constructors. This handout explores implicit type conversions, conversion constructors, and how to prevent coding errors stemming from inadvertent conversion constructors.

Implicit Conversions

In C++, an *implicit conversion* is a conversion from one type to another that doesn't require an explicit typecast. Perhaps the simplest example is the following conversion from an `int` to a `double`:

```
double myDouble = 137 + 2.71828;
```

Here, even though 137 is an `int` while 2.71828 is a `double`, C++ will implicitly convert it to a `double` so the operation can proceed smoothly.

When C++ performs implicit conversions, it does not “magically” figure out how to transform one data type into another. Rather, it creates a temporary object of the correct type that's initialized to the value of the implicitly converted object. Thus the above code is equivalent to

```
double temp = (double)myInt;  
double myDouble = temp + 2.71828;
```

It's important to remember that when using implicit conversions you are creating temporary objects. With primitive types this is hardly noticeable, but makes a difference when working with classes. For example, consider the following code:

```
string myString = "This ";  
string myOtherString = myString + "is a string";
```

Note that in the second line, we're adding a C++ `string` to a C `char *` string. Thus C++ will implicitly convert “is a string” into a C++ `string` by storing it in a temporary object. The above code, therefore, is equivalent to

```
string myString = "This ";  
string tempStr = "is a string";  
string myOtherString = myString + tempStr;
```

Notice that in both of the above examples, at some point C++ needed a way to initialize a temporary object to be equal to an existing object of a different type. In the first example, we made a temporary `double` that was equal to an `int`, and in the second, a temporary `string` equal to a `char *`. When C++ performs these conversions, it uses a special function called a *conversion constructor* to initialize the new object. Conversion constructors are simply class constructors that accept a single parameter and initialize the new object to a semantically-equivalent copy of the parameter. In the `double` example, the newly-created `double` had the same value as the `int` parameter. With the C++ `string`, the temporary `string` was equivalent to the C `string`.

C++ will invoke conversion constructors in all of the same places where it invokes copy constructors.* Thus, if you pass a `char *` to a function accepting a C++ `string`, the `string` will be initialized to the `char *` in its conversion constructor. Similarly, if you have a function like this one:

```
string MyFunction()
{
    return "This is a string!";
}
```

The temporary object created for the return value will be initialized to the C `string` "This is a string!" using the conversion constructor.

Writing Conversion Constructors

To see how to write conversion constructors, we'll use the example of a `CString` class that's essentially our own version of the C++ `string` class. Internally, `CString` stores the string as a C `string` called `theString`. Since we'd like to define an implicit conversion from `char *` to `CString`, we'll declare a conversion constructor, as shown below:

```
class CString
{
public:
    CString(const char *other);
    /* Other member functions. */
private:
    char *theString;
};
```

Then we'd implement the conversion constructor as

```
CString::CString(const char *other)
{
    /* Allocate space and copy over the string. */
    theString = new char[strlen(other) + 1];
    strcpy(theString, other);
}
```

Now, whenever we have a `char *` C `string`, we can implicitly convert it to a `CString`.

In the above case, we defined an implicit conversion from `char *` C `strings` to our special class `CString`. However, it's possible to define a second conversion from a C++ `string` to our new `CString` class. In fact, C++ allows you to provide conversion constructors for any number of different

* In fact, technically speaking, a copy constructor is a special case of the conversion constructor.

types that may or may not be primitive types.

Here's a modified `CString` interface that provides a copy constructor and two conversion constructors from `string` and `char *`:

```
class CString
{
    public:
        CString(const string &other);
        CString(const char *other);

        /* The rule-of-three says we need these functions. */
        CString(const CString &other);
        CString &operator =(const CString &other);
        ~CString();
        /* Other member functions. */
    private:
        char *theString;
};
```

A Word on Readability

When designing classes with conversion constructors, it's easy to get carried away by adding too many implicit conversions. For example, suppose that for the `CString` class we want to define a conversion constructor that converts `ints` to their string representations. This is completely legal, but can result in some confusing or unreadable code. For example, if there's an implicit conversion from `ints` to `CStrings`, then we can write code like this:

```
CString myStr = myInt + 137;
```

The resulting `CString` would then hold a string version of the value of `myInt + 137`, not the string composed of the concatenation of the value of `myInt` and the string “137.” This can be a bit confusing and can lead to counterintuitive code. Worse, since C++ does not normally define implicit conversions between numeric and string types, people unfamiliar with the `CString` implementation might get confused by lines assigning `ints` to `CStrings`.

In general, when working with conversion constructors, make sure that the conversion is intuitive and consistent with major C++ conventions. If not, consider using non-constructor member functions. For example, if we would like `CString` to be able to convert `int` values into their string representations, we might want to make a member function `intToString` that performs the conversion. This way, someone reading the code could explicitly see that we're converting an `int` to a `CString`.

Conversion Assignment Operators

One question that might have crossed your mind is whether there's a need for a “conversion assignment operator” to pair with the conversion constructor. After all, with our `CString` class, since we've only provided a single assignment operator, it would seem that we shouldn't be able to assign a regular `char *` to a `CString`.

It turns out, however, that the following code is legal:

```
CString myStr;  
myStr = "This is a C string!"; // Invokes assignment operator
```

The reason that this code is legal has to do with how C++ interprets the statement `myStr = "This is a C string!"`. Note that in our `CString` class we provided an assignment operator as the function `CString &operator =(const CString &other)`. This means that when we write the line `myStr = "This is a C string!"`, we're actually calling the `operator =` function. Thus the above code snippet is equivalent to

```
CString myStr;  
myStr.operator =("This is a C string!");
```

We'll go into code like this in more detail next week when discussing operator overloading, but for now just be aware that the above two lines are completely legal C++ code.

If you'll recall, since we're passing the value “This is a C string!” to a function that accepts an object of type `CString`, C++ will automatically invoke the `CString` conversion constructor to fill in the `CString` parameter. The above code, therefore, is equivalent to

```
CString myStr;  
CString tempStr("This is a C string!");  
myStr = tempStr;
```

And the assignment will work out correctly. In short, since we defined conversion constructors to convert `char *s` into `CStrings`, we're free to assign `char *s` to `CStrings` without defining a new assignment operator.

Problems with Conversion Constructors

While conversion constructors are quite useful in a wide number of circumstances, the fact that C++ automatically treats all single-parameter constructors as conversion constructors can lead to convoluted or nonsensical code.

One of my favorite examples of “conversion-constructors-gone-wrong” comes from an older version of the CS106 ADT class libraries. Originally, the CS106 `Vector` was defined as

```
template<typename ElemType>  
class Vector  
{  
    public:  
        Vector(int sizeHint = 10); // Hint about the size of the Vector  
        /* ... */  
};
```

Nothing seems all that out-of-the-ordinary here – we have a `Vector` template class that lets you give the class a hint about the number of elements you will be storing in it. However, because the constructor accepts a single parameter, C++ will interpret it as a conversion constructor and thus will let us implicitly convert from `ints` to `Vectors`. This can lead to some very strange behavior. For example, given the above class definition, consider the following code:

```
Vector<int> myVector = 15;
```

This code, while nonsensical, is legal and equivalent to `Vector<int> myVector(15)`. Fortunately, this probably won't cause any problems at runtime – it just doesn't make sense in code.

However, suppose we have the following code:

```
void DoSomething(Vector<int> &myVector)
{
    myVector = 0;
}
```

This code is totally legal even though it makes no logical sense. Since the code expands to a call to `myVector.operator=(0)`, C++ will create a new `Vector<int>` initialized with the parameter 0 and then assign it to `myVector`. In other words, the above code is equivalent to

```
void DoSomething(Vector<int> &myVector)
{
    Vector<int> tempVector(0);
    myVector = tempVector;
}
```

`tempVector` is empty when it's created, so when we assign `tempVector` to `myVector`, we'll set `myVector` to the empty vector. Thus the nonsensical line `myVector = 0` is effectively an obfuscated call to `myVector.clear()`.

This is a quintessential example of why conversion constructors can be dangerous. When writing single-argument constructors, you run the risk of letting C++ interpret your constructor as a conversion constructor.

explicit

To prevent problems like the one described above, C++ provides the `explicit` keyword to indicate that a constructor cannot be interpreted as a conversion constructor.

For example, let's look at the current version of the CS106 `Vector`, which has its constructor marked `explicit`:

```
template<typename ElemType>
class Vector
{
public:
    explicit Vector(int sizeHint = 10); // Hint the size of the Vector
    /* ... */
};
```

Now, if we write code like

```
Vector<int> myVector = 10;
```

We'll get a compile-time error since there's no implicit conversion from `int` to `Vector<int>`. However, we can still write

```
Vector<int> myVector(10);
```

Which is what we were trying to accomplish in the first place. Similarly, we eliminate the `myVector = 0` error, and a whole host of other nasty problems.

When designing classes, if you have a single-argument constructor that is not intended as a conversion function, you *must* mark it `explicit` to avoid running into the “implicit conversion” trap. While indeed this is more work for you as an implementer, it will make your code safer and more stable.

More to Explore

If you're interested in writing your own container classes, you might want to consider reading into *generalized copy constructors*. A generalized copy constructor is a templated copy constructor that accepts objects of the same template class whose template types don't match the current object's. For example, given a `Vector` template, you might define a generalized copy constructor so that you could initialize a `Vector<double>` to the contents of a `Vector<int>` or vice-versa. Of course, the syntax for generalized copy constructors can be a bit tricky (it requires the `friend` keyword, which we'll cover next week), but such is the nature of heavily-templated code!

Practice Problems

These practice problems concern a `RationalNumber` class that encapsulates a rational number (that is, a number expressible as the quotient of two integers). `RationalNumber` is declared as follows:

```
class RationalNumber
{
    public:
        RationalNumber(int num = 0, int denom = 1) :
            numerator(num), denominator(denom) {}

        double getValue() const {
            return double(numerator) / denominator;
        }

        void setNumerator(int value) {
            numerator = value;
        }
        void setDenominator(int value) {
            denominator = value;
        }
    private:
        int numerator, denominator;
};
```

The constructor to `RationalNumber` accepts two parameters that have default values. This means that if you omit one or more of the parameters to `RationalNumber`, they'll be filled in using the defaults. Thus all three of the following lines of code are legal:

```
RationalNumber zero; // Value is 0 / 1 = 0
RationalNumber five(5); // Value is 5 / 1 = 5
RationalNumber piApprox(355, 113); // Value is 355/113 = 3.1415929203...
```

1. Explain why the `RationalNumber` constructor is a conversion constructor.
2. Write a `RealNumber` class that encapsulates a real number (any number on the number line). It should have a conversion constructor that accepts a `double` and a default constructor that sets the value to zero. (*Note: You only need to write one constructor. Use `RationalNumber` as an example*)
3. Write a conversion constructor for `RealNumber` that converts a `RationalNumber` to a `RealNumber`.

General questions:

5. If a constructor has two or more arguments and no default values, can it be a conversion constructor?
6. C++ will apply at most one implicit type conversion at a time. That is, if you define three types `A`, `B`, and `C` such that `A` is implicitly convertible to `B` and `B` is implicitly convertible to `C`, C++ will not automatically convert objects of type `A` to objects of type `C`. Give an reason for why this might be. (*Hint: Add another implicit conversion between these types*)
7. What happens when you mark a copy constructor `explicit`? (*Hint: Will the syntax `MyClass one = two` be legal? How about `MyClass one(two)?`*)
8. Write a conversion constructor converting a C++ string to our special `CString` class.